

https://i0.wp.com/thedatascientist.com/wp-content/uploads/2018/05/recommender_systems.png

Lecture 5: Scalable Matrix Factorisation for Collaborative Filtering in RecSys

[COM6012: Scalable ML](#) by [Haiping Lu](#)


YouTube Playlist: <https://www.youtube.com/c/HaipingLu/>

Week 5 Contents / Objectives

- Recommender Systems
- Collaborative Filtering
- Matrix Factorisation for Collaborative Filtering
- Scalable Collaborative Filtering in Spark

Week 5 Contents / Objectives

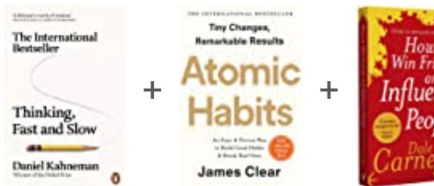
- Recommender Systems
- Collaborative Filtering
- Matrix Factorisation for Collaborative Filtering
- Scalable Collaborative Filtering in Spark



<https://blogs-image>


Recommendations Everywhere

Frequently bought together



- ✓ **This item:** Thinking, Fast and Slow by D
- ✓ **Atomic Habits:** The life-changing million
- ✓ **How to Win Friends and Influence People**

Customers who viewed this



Nudge: Improving Decisions About Health, Wealth and Happiness
› Richard H. Thaler

Predictable Patterns: The Hidden Forces That Shape Our Lives
› Dan A

Google

job

- jobs at amazon uk
- job in sheffield
- jobs near me
- jobs sheffield
- jobs ecclesall road
- jobs
- jobs ac uk
- jobs at amazon
- Jobcentre Plus
- jobs indeed uk

Google Search

Online events for you

See all



Scholars Webinar on Drug Discovery...

Wed, Mar 24 - Thu, M...

Alireza Khanteymoori and 1,616 other...

View



Accelerating Visual Data Exploration...

Today, 12:30 PM

96 attendees

View



Live Chat with Salesforce Sr...

Tue, May 11, 5:30 PM

Vignesh Srinivasan and 733 other attendees

View

People you may know from The University of Sheffield

See all



Twin Karma...

Research Software...
9 mutual connections

Connect



Andrew Stra...

Senior University Teacher...
19 mutual connections

Connect



Neil Walkins...

Senior Lecturer at The Universi...
30 mutual connections

Connect



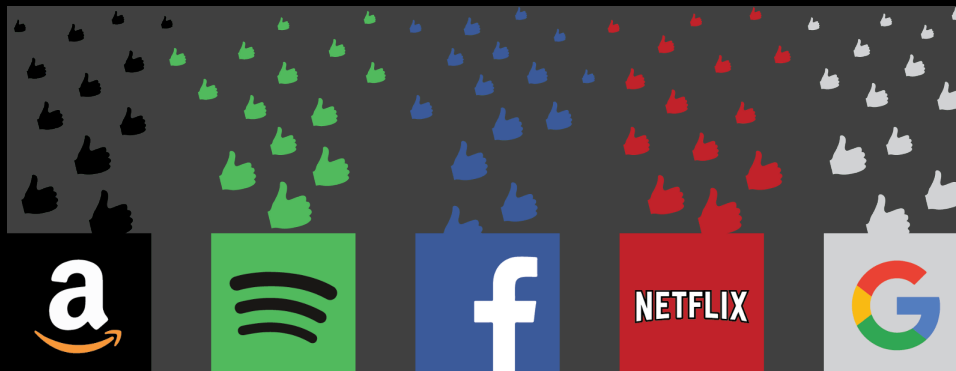
Siobhan No...

Senior Lecturer at University of...
21 mutual connections

Connect

Recommender Systems (RecSys)

- **Predict** relevant items for a user, in a given context
- **Predict** to what extent these items are relevant
- A **ranking** task (**searching** as well)
- Implicit, targeted, **intelligent** advertisement
- Effective, popular marketing

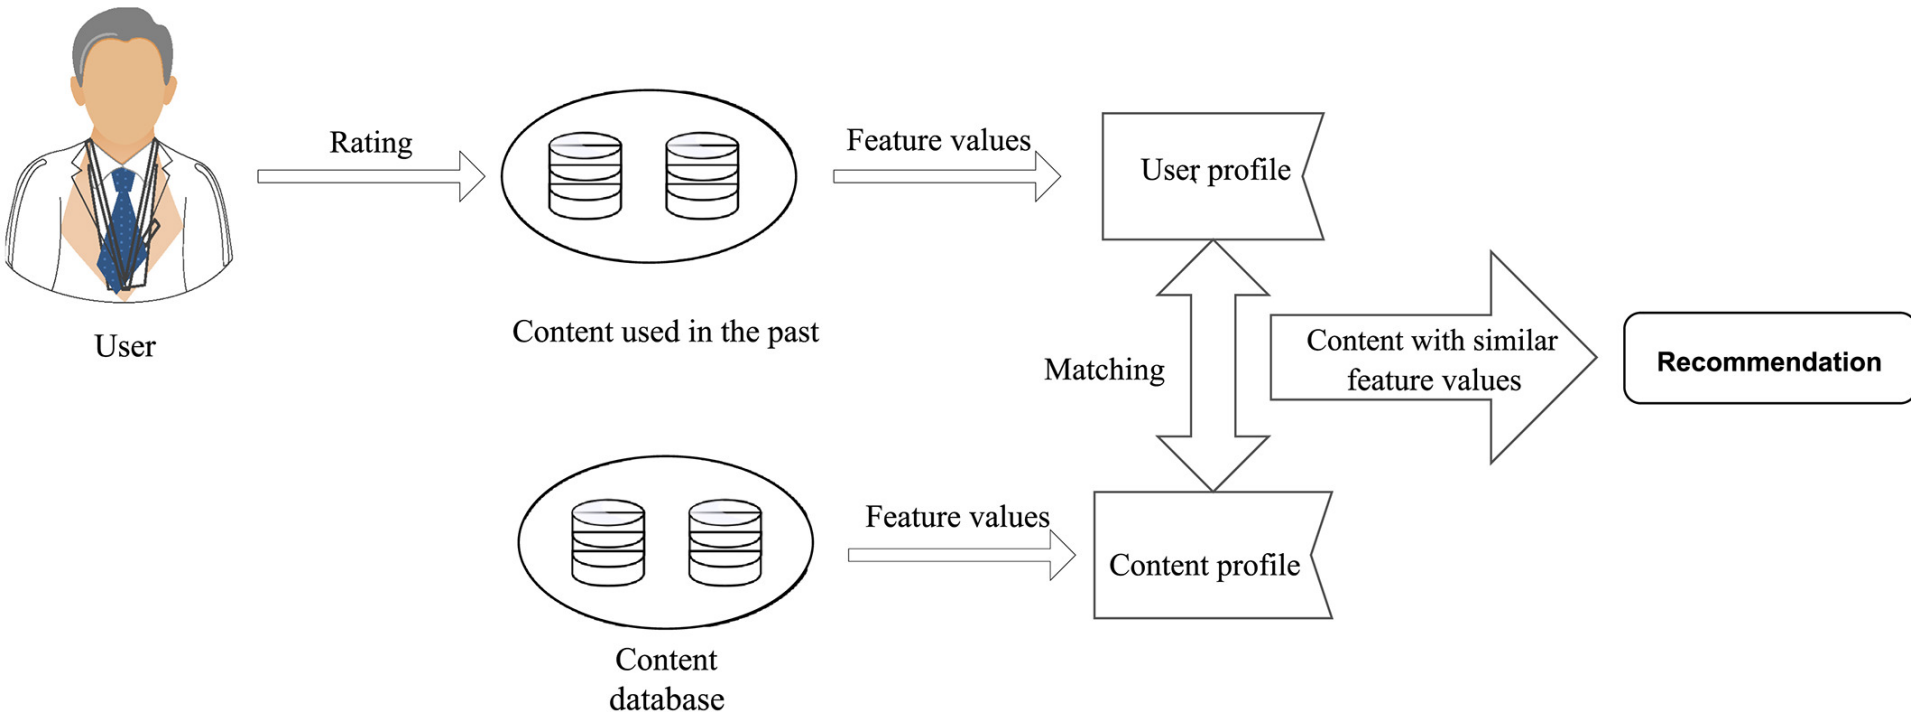


Two Classes of RecSys

- **Content-based** recommender systems
- **Collaborative filtering** recommender systems

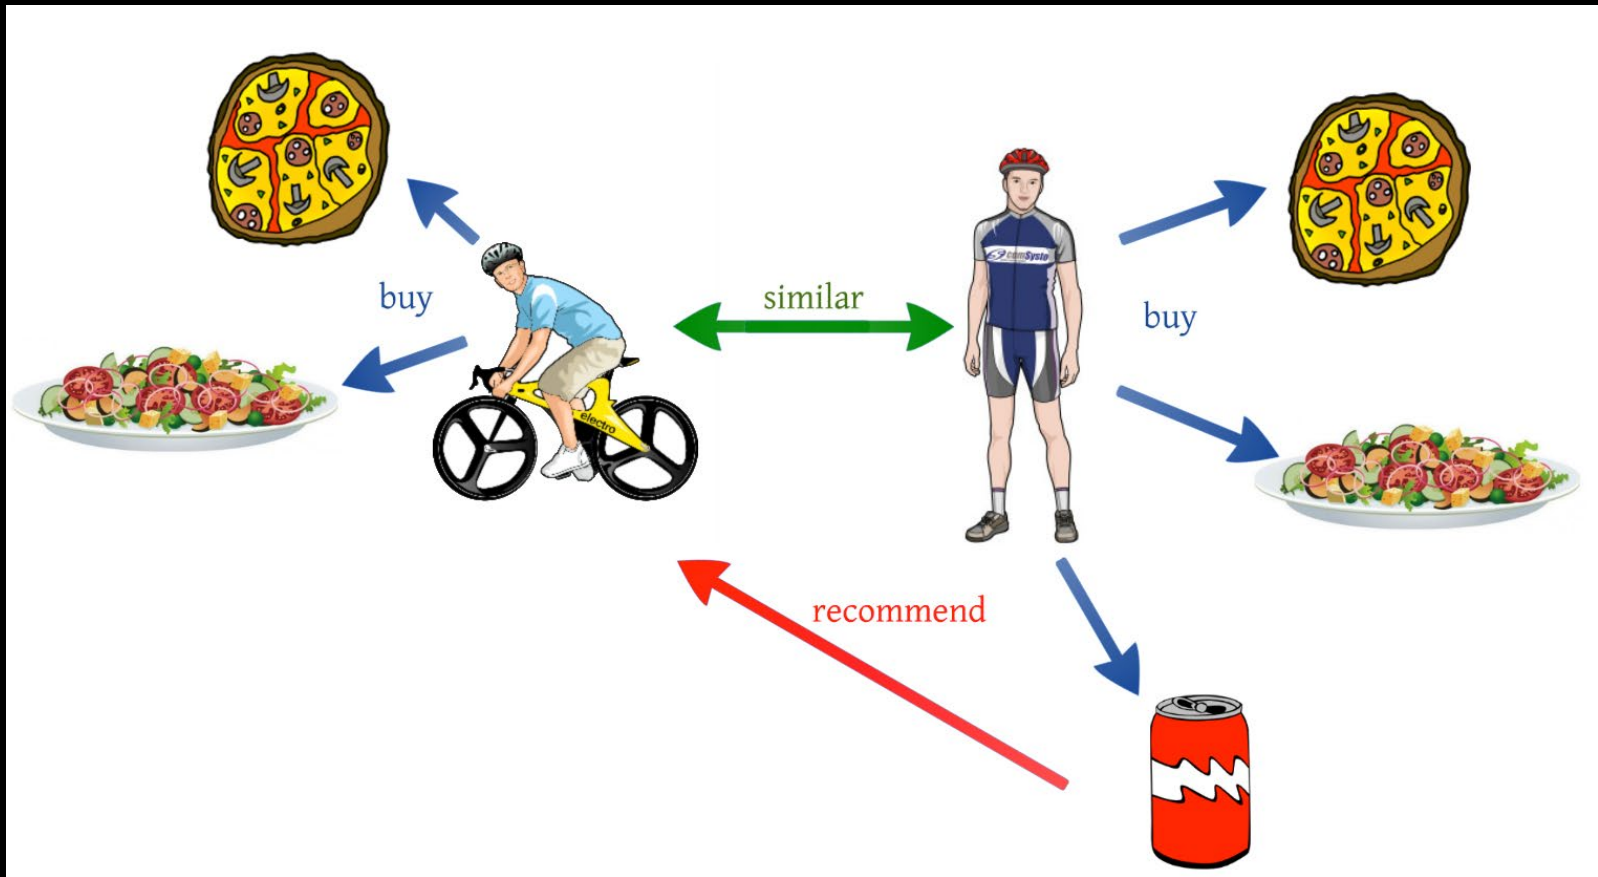


Content-based RecSys



[dac4058-fig-0010-m.jpg \(2128×789\) \(wiley.com\)](#)

Collaborative Filtering RecSys




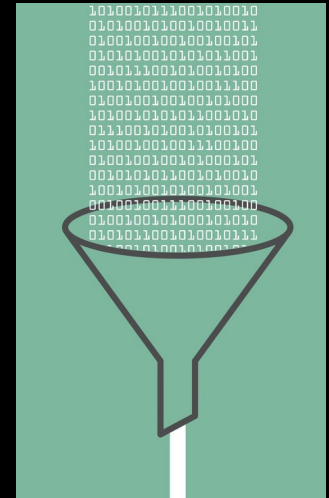
https://miro.medium.com/max/2656/1*6_NIX6CJYhtxzRM-t6ywkQ.png

Week 5 Contents / Objectives

- Recommender Systems
- Collaborative Filtering
- Matrix Factorisation for Collaborative Filtering
- Scalable Collaborative Filtering in Spark

What is Collaborative Filtering?


- **Information** filtering based on past records
 - Electronic word of mouth marketing
 - Turn visitors into customers (e-Salesman)
- Components
 - **Users** (customers): who provide ratings
 - **Items** (products): to be rated
 - **Ratings** (interest): core data



				
John	5	1	3	5
Tom	?	?	?	2
Alice	4	?	3	?


Collaborative Filtering (CF)

- Objective: predict how well a user will like an **unrated** item, given past ratings for a community of users
- How does CF work?
 - Input: many users' **ratings** for many items
 - Model: similar users \leftarrow ratings strongly correlate
 - Recommend items rated highly by similar users



Explicit vs Implicit Ratings

- Explicit (direct): users indicate levels of interest
 - Most accurate descriptions of a user's preference
 - Challenging in collecting data
- Implicit (indirect): observing user behavior
 - Can be collected with little or no cost to user
 - Ratings inference may be imprecise



Rating Scales

- Scalar ratings

- Numerical scales
- 1-5, 1-7, etc.



- Binary ratings

- Agree/Disagree, Good/Bad, etc.



- Unary ratings

- Presence/absence of an event, e.g., purchase/browsing history, search patterns, mouse movements
- No info about the opposite $\neq 0$



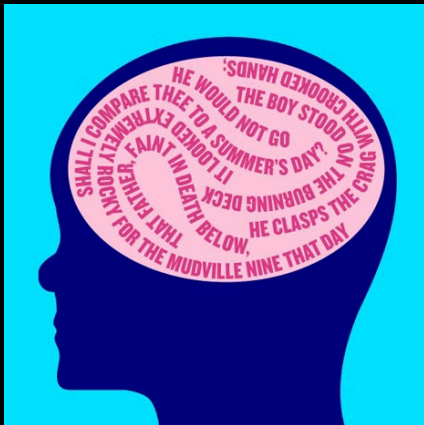
CF Preferences

- Many users, many items, many ratings
- Users rate multiple items
- Other users with similar needs/tastes
- Item evaluation requires personal taste
- Taste persists



CF Methods

- Memory-based: predict using past ratings **directly**
 - Weighted ratings given by other similar users
 - User-based & item-based (non-ML)
- Model-based: **model** users based on past ratings
 - Predict ratings using the learned model



[iforget-465.jpg \(465×465\) \(newyorker.com\)](#)



[neural-header.jpg \(756×503\) \(utsouthwestern.edu\)](#)

Prediction Accuracy

- Mean absolute error (MAE)

$$MAE = \frac{\sum_{i,j} |p_{i,j} - r_{i,j}|}{n}$$

- Normalized MAE

$$NMAE = \frac{MAE}{r_{max} - r_{min}}$$

- Root mean squared error (RMSE)

$$RMSE = \sqrt{\frac{1}{n} \sum_{i,j} (p_{i,j} - r_{i,j})^2}$$

Challenges

- **Cold Start**

- New user

- Rate some initial items
 - Non-personalized rec.
 - Describe tastes
 - Demographic info

- New item


- Randomly selecting items
 - Content analysis, metadata (non-CF)

- **Sparsity**: sparse user-item matrix

- **Scalability**: millions of users and items



[isbil+2.jpg \(490×303\) \(squarespace-cdn.com\)](#)



Week 5 Contents / Objectives

- Recommender Systems
- Collaborative Filtering
- **Matrix Factorisation for Collaborative Filtering**
- Scalable Collaborative Filtering in Spark

Matrix Factorisation (MF) for CF

- Characterise items/users by vectors of factors learned from the rating matrix **user x item**
- High correlation between item and user factors → good recommendation
- Flexibility: incorporate implicit feedback, temporal effects, and confidence levels

				
John 	5	1	3	5
Tom 	?	?	?	2
Alice 	4	?	3	?


Basic MF Model

- Map users & items to a **joint latent factor** space of dimensionality k
 - Item $i \rightarrow$ vector q_i : the extent to which the item possesses those k factors
 - User u : vector p_u : the extent of interest the user has on those k factors
- User-item interactions: the user's overall interest in the item's characteristics
 - Inner product $q_i^T p_u$: predicted user u 's rating of item i

$$\hat{r}_{u,i} = q_i^T p_u$$

How to Learn the MF Model



- To learn: item factors $\{q_i\}$ and user factors $\{p_u\}$
 - Factorisation assuming full rating matrix
 - Factorise rating matrix R using SVD to obtain P, S, Q
- $$R = PSQ^T$$
- Reduce the matrix S to dimension k , i.e. S_k
 - $P \rightarrow P_k$ and $Q \rightarrow Q_k: P_k S_k \rightarrow \hat{P}$, and $S_k Q_k^T \rightarrow \hat{Q}^T$
 - u th row of $\hat{P} \rightarrow p_u$, i th column of $\hat{Q}^T \rightarrow q_i$



$A: R$
 $U: P$
 $\Sigma: S$
 $V: Q$
 $r: k$


Challenges in MF for CF

- High portion of missing values caused by sparseness in the user-item rating matrix
- Conventional SVD is undefined when knowledge about the matrix is incomplete



How to Fill Missing Values

- Imputation: fill in missing ratings using the average ratings for user and item
- Problems
 - Expensive: significantly increases the amount of data
 - Inaccurate imputation might distort the data



MF with Missing Values

- Modelling directly the **observed ratings only**
 - Avoid overfitting through a regularised model
 - Minimize the regularised squared error on the set of known ratings to learn the factor vectors p_u and q_i

$$\min_{q^*, p^*} \sum_{(u,i) \in \kappa} (r_{ui} - q_i^T p_u)^2 + \lambda (\|q_i\|^2 + \|p_u\|^2)$$

- κ : the training set of the (u,i) pairs with known ratings
- λ : the regularisation parameter

Alternating Least Squares for MF-CF

$$\min_{q^*, p^*} \sum_{(u,i) \in \kappa} (r_{ui} - q_i^T p_u)^2 + \lambda(\|q_i\|^2 + \|p_u\|^2)$$

- Both p_u and q_i are unknown (non-convex function)
- Fix one of them \rightarrow quadratic with optimal solution
- Alternating Least Squares (ALS): alternate between fixing q_i s and fixing p_u s
 - Fix P (p_u s) as \hat{P} to recompute q_i s by solving a least-squares problem $\|R - PQ^T\|_F$ ([Frobenius norm](#))

$$R = \hat{P}Q^T \Rightarrow Q^T = (\hat{P}^T \hat{P})^{-1} \hat{P}^T R$$

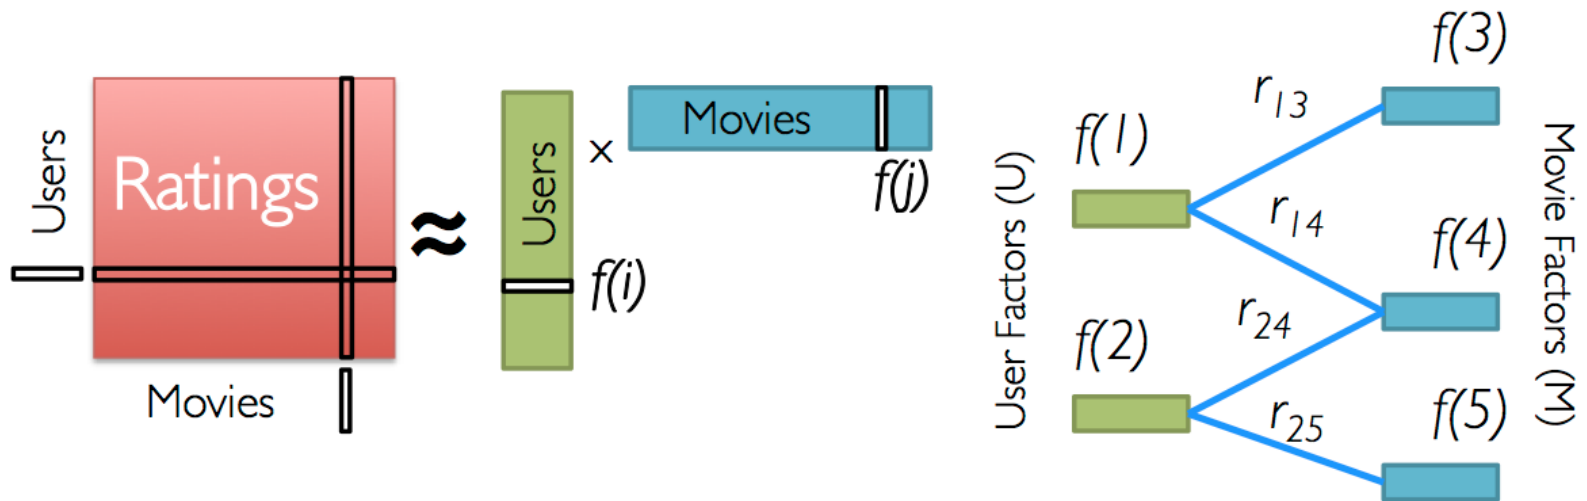
- Fix Q as \hat{Q} , we have

$$P = R\hat{Q}(\hat{Q}^T \hat{Q})^{-1}$$

- Random initialisation to start this iteration

MF for Movie Recommendation

Low-Rank Matrix Factorization:



Iterate:

$$f[i] = \arg \min_{w \in \mathbb{R}^d} \sum_{j \in \text{Nbrs}(i)} (r_{ij} - w^T f[j])^2 + \lambda ||w||_2^2$$

[spark-training/matrix_factorization.png](https://github.com/databricks/spark-training/blob/master/matrix_factorization.png) at master · databricks/spark-training (github.com)

Week 5 Contents / Objectives


- Recommender Systems
- Collaborative Filtering
- Matrix Factorisation for Collaborative Filtering
- Scalable Collaborative Filtering in Spark

Key in Scalable ML

- Computation and storage should be **linear** (in n, d)
→ Low-cost computation (time + space)
- Perform **parallel** and **in-memory** computation
→ Many working + reduce disk I/O
- Minimise network **communication** → Reduce overhead in parallelisation, not the more the better

Blocked Implementation of ALS

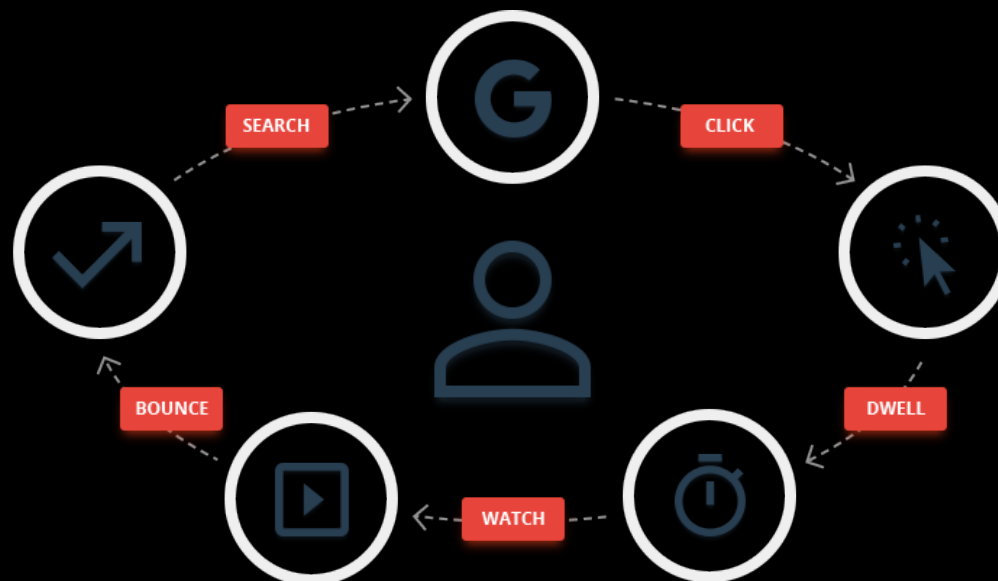
- Group users and items into blocks
 - Reduce **communication**: only send one copy of each user vector to each item block on each iteration, and only for the item blocks that need that user's feature vector
 - **Pre-compute** info: **out-links** of each user (which blocks of items it will contribute to); **in-links** for each item (which of the feature vectors it receives from each user block it will depend on)



[Color-online-A-symmetric-block-Toeplitz-matrix-Each-block-is-also-a-symmetric-Toeplitz.png \(488x369\) \(researchgate.net\)](#)

Implicit Feedback Modelling

- Implicit feedback: views, clicks, purchases, likes, shares
 - Rating r = strength in observations of user actions (#clicks, viewing duration) \rightarrow confidence level in observed user preference
 - Construct a preference matrix P : e.g. 1 if $r > 0$ and 0 if $r = 0$
 - Factorisation of $P \rightarrow$ latent factors to predict the preference of a user for an item (details in an [ICDM08 paper](#))



The ALS API in Spark

- **numUserBlocks/numItemBlocks**: the number of blocks the users/items will be partitioned into to parallelize computation (defaults to 10)
- **rank**: the number of latent factors in the model (defaults to 10)
- **regParam**: the regularization parameter in ALS (defaults to 0.1)
- **implicitPrefs**: whether to use the explicit feedback ALS variant or one adapted for implicit feedback data (defaults to false: explicit ratings)
- **alpha**: the baseline confidence of implicit feedback (defaults to 1.0)
- **nonnegative**: whether to use nonnegative constraints (defaults to false)
- **coldStartStrategy**: “drop” → drop any rows in the DataFrame of predictions that contain NaN values (defaults to “nan”: assign NaN to a user and/or item factor is not present in the model)
- **blockSize**: the size of the user/product blocks in the blocked implementation of ALS to reduce communication

```

944 def train[ID: ClassTag]( // scalastyle:ignore
945     ratings: RDD[Rating[ID]],
946     rank: Int = 10,
947     numUserBlocks: Int = 10,
948     numItemBlocks: Int = 10,
949     maxIter: Int = 10,
950     regParam: Double = 0.1,
951     implicitPrefs: Boolean = false,
952     alpha: Double = 1.0,
953     nonnegative: Boolean = false,
954     intermediateRDDStorageLevel: StorageLevel = StorageLevel.MEMORY_AND_DISK,
955     finalRDDStorageLevel: StorageLevel = StorageLevel.MEMORY_AND_DISK,
956     checkpointInterval: Int = 10,
957     seed: Long = 0L)(
958     implicit ord: Ordering[ID]): (RDD[(ID, Array[Float])], RDD[(ID, Array[Float])]) = {
959
960     require(!ratings.isEmpty(), s"No ratings available from $ratings")
961     require(intermediateRDDStorageLevel != StorageLevel.NONE,
962         "ALS is not designed to run without persisting intermediate RDDs.")
963
964     val sc = ratings.sparkContext
965
966     // Precompute the rating dependencies of each partition
967     val userPart = new ALSPartitioner(numUserBlocks)
968     val itemPart = new ALSPartitioner(numItemBlocks)
969     val blockRatings = partitionRatings(ratings, userPart, itemPart)
970         .persist(intermediateRDDStorageLevel)
971     val (userInBlocks, userOutBlocks) =
972         makeBlocks("user", blockRatings, userPart, itemPart, intermediateRDDStorageLevel)
973     userOutBlocks.count() // materialize blockRatings and user blocks

```

CF in Spark ML

- [Scala code](#) (1800+ lines)
- Documentation: [Collaborative Filtering in Spark](#)
- [DataBricks movie recommendations tutorial](#)
- [DataBricks](#): founded by the creators of Apache Spark
 - Their latest packages at [their GitHub page](#)
 - [Databricks community edition](#): 14 days free on AWS, Microsoft Azure or Google Cloud.



References

- Yehuda Koren, Robert Bell, and Chris Volinsky.
"[Matrix factorization techniques for recommender systems.](#)" *Computer* 8 (2009): 30-37 (Yahoo & AT&T)
- Yifan Hu, Yehuda Koren, and Chris Volinsky.
"[Collaborative filtering for implicit feedback datasets.](#)" *Eighth IEEE International Conference on Data Mining*, 2008
- Charu C. Aggarwal, [Recommender Systems: The Textbook](#), April 2016